

Detection of Canned Beverages

1st Radu Ștefan-Vlăduț

Cohort 1306A

stefan-vladut.radu@student.tuiasi.ro

2nd Petrișor Eduard-Gabriel

Cohort 1311A

eduard-gabriel.petrisor@student.tuiasi.ro

I. INTRODUCTION

With the given task of building a program with the function of detecting certain objects by using an individually procured data set we initially pondered what thing we might tailor our project towards, considering the vastness of the things we could've chosen. After much deliberating, we set our sights on the humble and common soda can. With different shapes, sizes and colors, it offers plenty of difference from one brand to another whilst still presenting a unifying form, all that on top of being a readily available object which can be found at any corner store, which greatly aids the task of making our very own data set. But what truly swayed us towards this idea is the difficulty of actually distinguishing the different kinds of cans by use of only your tactile senses. Whilst some cans might present different textures and reliefs, some are essentially virtually indistinguishable from one another, fact which we hope offers our project a kind of real world use.

II. STATE-OF-THE-ART

Object detection performance, as measured on the canonical PASCAL VOC dataset, has plateaued in the last few years. The best-performing methods are complex ensemble systems that typically combine multiple low-level image features with high-level context. The state-of-the-art methods can be categorized into two main types: one-stage methods and two stage-methods:

A. One-stage methods

One-stage methods prioritize inference speed, and example models include YOLO, SSD and RetinaNet.

- **YOLO:** YOLO is a state-of-the-art, real-time object detection system that can detect over 9000 object categories. YOLO suffers from a variety of shortcomings relative to state-of-the-art detection systems. Error analysis of YOLO compared to Fast R-CNN shows that YOLO makes a significant number of localization errors. Furthermore, YOLO has relatively low recall compared to region proposal-based methods. Thus it focuses mainly on improving recall and localization while maintaining classification accuracy. It's main defining quality however is being extremely easy and out of the box easy to use.

- **SSD:** The SSD approach is based on a feed-forward convolutional network that produces a fixed-size collection of bounding boxes and scores for the presence of object class instances in those boxes, followed by a non-maximum suppression step to produce the final detections. Is the first deep network based object detector that does not resample pixels or features for bounding box hypotheses and is as accurate as approaches that do. This results in a significant improvement in speed for high-accuracy detection. The fundamental improvement in speed comes from eliminating bounding box proposals and the subsequent pixel or feature resampling stage. Improvements include using a small convolutional filter to predict object categories and offsets in bounding box locations, using separate predictors (filters) for different aspect ratio detections, and applying these filters to multiple feature maps from the later stages of a network in order to perform detection at multiple scales.
- **RetinaNet:** Similar to SSD, RetinaNet uses Feature Pyramid Networks (FPN) to extract multiple-scale information and performs classification and box regression in one stage. In RetinaNet, a novel loss function is adopted to address class foreground and background class extremely imbalance. With such a novel loss function, RetinaNet can achieve high accuracy in object detection with fast inference time. RetinaNet can use different neural networks as the backbone model to learn feature maps, which are the input to the FPN. The third component of RetinaNet is the classification subnet and box regression subnet attached to FPN to compute the results. For example, the max-pooling output of last three blocks in VGG can be chosen as FPN input in RetinaNet, while the output of last layer in Inception-ResNet block can be selected as the FPN input.

B. Two-stage methods

Two-stage methods prioritize detection accuracy, and example models include Faster R-CNN, Mask R-CNN and Cascade R-CNN.

- **Faster R-CNN:** Faster R-CNN, is composed of two modules. The first module is a deep fully convolutional network that proposes regions, and the second module is the Fast R-CNN detector [2] that uses the proposed regions. The entire system is a single, unified network for object detection. Using the recently popular terminology of neural networks with 'attention' mechanisms, the RPN

module tells the Fast R-CNN module where to look. A Region Proposal Network (RPN) takes an image (of any size) as input and outputs a set of rectangular object proposals, each with an objectness score.³ We model this process with a fully convolutional network, which we describe in this section. Because our ultimate goal is to share computation with a Fast R-CNN object detection network, we assume that both nets share a common set of convolutional layers. In our experiments, we investigate the Zeiler and Fergus model (ZF), which has 5 shareable convolutional layers and the Simonyan and Zisserman model (VGG-16), which has 13 shareable convolutional layers.

- **Mask R-CNN:** Mask R-CNN, extends Faster R-CNN by adding a branch for predicting segmentation masks on each Region of Interest (RoI), in parallel with the existing branch for classification and bounding box regression (Figure 1). The mask branch is a small FCN applied to each RoI, predicting a segmentation mask in a pixel-to-pixel manner. Mask R-CNN is simple to implement and train given the Faster R-CNN framework, which facilitates a wide range of flexible architecture designs. Additionally, the mask branch only adds a small computational overhead, enabling a fast system and rapid experimentation.
- **Cascade R-CNN:** It is a multi-stage extension of the R-CNN, where detector stages deeper into the cascade are sequentially more selective against close false positives. The cascade of R-CNN stages are trained sequentially, using the output of one stage to train the next. This is motivated by the observation that the output IoU of a regressor is almost invariably better than the input IoU, where nearly all plots are above the gray line. It suggests that the output of a detector trained with a certain IoU threshold is a good distribution to train the detector of the next higher IoU threshold. This is similar to bootstrapping methods commonly used to assemble datasets in object detection literature. The main difference is that the resampling procedure of the Cascade R-CNN does not aim to mine hard negatives. Instead, by adjusting bounding boxes, each stage aims to find a good set of close false positives for training the next stage. When operating in this manner, a sequence of detectors adapted to increasingly higher IoUs can beat the overfitting problem, and thus be effectively trained. At inference, the same cascade procedure is applied. The progressively improved hypotheses are better matched to the increasing detector quality at each stage. This enables higher detection accuracies.

After heavily surveying the above mentioned documents and weighing our options we decided to use YOLO for our project due to it being fairly simple and easy to use out of the box as well as the widely prevalent resources for it.

RELATED WORK

We have combed through different research documents and implementations regarding this subject and came to the initial question of if we would try using a one-stage or a two-stage implementation. Through the study of reference [7] we learned the characteristics of the R-CNN derived two-stage implementations and whilst they prove to be superior to the one-stage implementations in most metrics, they seemed to be far too complex for the scope of our project. As such, we decided on a one-stage method and after we looked through materials [4] and [5] we ended on using YOLO for our project.

METHOD DESCRIPTION

The main defining feature of our project compared to the others is that we had to procure our very own data set to train our YOLO-based network on. Whilst this did seem like an easy prospect as first, we soon realised the magnitude of the task. After acquiring the images we used a free online tool called CVAT (Computer vision annotation tool) to annotate every image individually. Annotating an image consists of drawing a bounding rectangle around the object you wish the algorithm to detect. With the annotating process done, we used another online tool called Roboflow to split our annotated images into 3 categories: training images, which would be used to train the AI itself and make it recognise the different objects, validation images, which are also used in the training process to ensure that the selected objects are correct and test images which serve the purpose of testing if the AI's training was successful by trying to verify the human placed bounding rectangles to the ones placed by the AI. With all of the data set fully formatted we can take to the actual coding phase of the project where the simplicity of the YOLO model comes into play. By importing our data from the Roboflow directly into our Colab Notebook with a few lines of code we can start training the model just by calling the 'train' function and simply specifying the data set it needs to use as well as the epochs, which entail the amount of times it should train itself on the data set. With all of that done, we have a functioning model that can be tested by importing it into another code file and giving it an image link alongside the "predict" command.

PRELIMINARY RESULTS AND CONCLUSIONS

At the moment we have a trained model that can fairly confidently detect cans with confidence of around 80-90 percent. On top of that we already have a bit of 200 images in our data set. One problem that we did run into however is that we initially wanted to make the model detect the individual flavors of the beverage, this however proved a herculean task with the relatively small size of the aforementioned data set, as such we had to relegate ourselves to just detect the brand as a whole. Biggest issue with this however was that we only had data for 4 of 6 types the brand has, making the model have a fairly difficult time recognising the 2 missing brands, which did hurt its prediction confidence quite a lot. A positive aspect we did notice however was the model's ability to recognise a can with

the same prediction no matter what side it was facing, this was mainly due to our decision to rotate the soda cans as we photographed them. A second issue that came with procuring our own data set was definitely that we did not factor in the background for all the photos which led to the program having severe problems when it was faced to complex backgrounds as well as being confused when faced with cans in awkward positions such as having them be laid on their side, on top of each other, at an angle or having them be very close one to another, to the point where the edges weren't clear. An issue that we were aware of and we tried our best to and hopefully did avoid is over fitting the AI, which happens when it is trained with the same images for far too many times. The conclusions we can draw from the current stage of the project is that the data set needs to be greatly enlarged and present a greater variation of positions for the objects as well as a far more varied backdrop for them.

REFERENCES

- [1] Mask R-CNN IEEE Paper
- [2] Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks
- [3] Cascade R-CNN: Delving Into High Quality Object Detection
- [4] SSD: Single Shot MultiBox Detector
- [5] YOLO9000 Better, Faster, Stronger
- [6] Road Damage Detection Using RetinaNet
- [7] Rich feature hierarchies for accurate object detection and semantic segmentation